⭐ POPULARITY

HUMAN Follow the crowd

AGENT Follow the corpus

48%

unnecessary NumPy usage

- Express picked for every JS backend
- npm over bun (Claude Code lead devs CLAUDE.md)
- Polars has 2x star growth; agents never pick it

Twist et al. (2025)

⏳ TEMPORAL

HUMAN Recency bias

AGENT Frozen at training cutoff

---

1 in 5

AI-recommended deps are safe

- javax.persistence instead of jakarta
- Ask for Intercom alternative, get Zendesk
- Training data circa 2022-2024 dominates

Endor Labs (2025)

⚓ ANCHORING

HUMAN First info dominates

AGENT Prior context bleeds across session

---

cross-domain

URL patterns misapplied

- mongodb.com/docs → python.org/docs
- shadcn/ui inserted unprompted into React
- Training prevalence + Tailwind synergy + anchoring

Huang et al. (2025)

🔄 DEFAULT

HUMAN Can't change established behavior

AGENT Can't change trained behavior

58%

Python for perf tasks. 0% Rust.

- Python chosen even when suboptimal
- Rust never selected across 8 models
- Override requires explicit context window rules

Twist et al. (2025)

SURVIVORSHIP

HUMAN Neglect the filtered-out

AGENT Invisible if not in corpus

2x

Polars star growth vs pandas; invisible to agents

- Plotly: 1 problem, 3 models
- Better library ≠ visible library
- Not in training → not in inference

Twist et al. (2025)

👻 CONFABULATION

HUMAN No clean human analogue

AGENT Invents missing libraries confidently

---

sparse

Pareto boundary: quality vs hallucination

- Hallucinated packages get registered as malware
- Slopsquatting: supply chain attack vector
- Agent is wrong with full conviction

Krishna et al. (2025)

Get found

through the stack.

Get preferred

through the corpus.